

GamerGain — Launch Cost Estimate (2026-07-30)

GamerGain — Launch Cost Estimate (floor edition)

☆ CURRENT RECONCILED ESTIMATE — 2026-08-15

All GamerGain / Get Goods Gratis launch-cost docs use these figures. Where an older number appears anywhere below, this banner supersedes it. (Reflects everything built through 2026-08-15: the AI concierge funnel + auto-launch, the public /Apply page marketing the Founding offer + “coming soon” financing, the results/benchmark auto-compilers and per-product stats, the Tier 2 “Scale” pay-as-you-go ladder + deliverables, the founding rollover/upgrade discount + \$2,000 premium gift boost, and the stronger cost floor — one-click in-app + `AI_FORCE_CHEAP_TIER` that dumps every AI call onto the small Llama. The /Apply page is a read + a lead form, so it adds capability, not launch or runtime cost.)

- **The launch number did NOT rise.** Every feature added this session ships **prebuilt and ON**, so it adds capability, not launch cost. The credit features (flexpay / financed Tier 1 / goods advance) ship **OFF by design** behind a counsel gate, so they add **\$0** legal cost at launch.
- **Full launch — web PWA + Android + native iOS: ~\$2,900-\$4,000 year-one all-in.** Developer labor ~\$2,050-\$3,150 (≈ 28-42 hours @ \$75/hr) — deploy, test, submit, and stand up the scheduler; no product build.
- **Shoestring — web PWA + Android (iOS as fast-follow): ~\$2,000-\$2,800.**
- **Hard external cash floor: ~\$139** all three platforms (Apple \$99/yr + Google Play \$25 + domain ~\$15); **~\$40 shoestring** (no Apple fee — the PWA covers iPhone via Safari).
- **Recurring: AI / media / email \$0/mo** — free tiers + the one-click cost floor; with `AI_FORCE_CHEAP_TIER` on, every call runs on Llama-8B so the LLM line is now ~\$0-20/mo (was \$5-40). Hosting ~\$10-35/mo (backend + Postgres + the always-on **scheduler** service). **Legal \$0** at launch (credit features off).
- **Standing auto-scaling is now ON from day one** (governor `AUTO_SCALE_ENABLED=1`, Railway provider, min 2 instances) — and the full-scale ceiling is now **hard-capped at ~\$3,000** by a codeable monthly-budget lever (`INFRA_SCALE_MONTHLY_BUDGET_USD`, default \$48/mo → ≤ 4 steady-state replicas ≈ \$576/yr), down from the old uncapped ~\$4,800 top end. A real surge still auto-bursts past the budget up to an 8-replica emergency ceiling (`INFRA_SCALE_BURST_ABOVE_BUDGET`, ON) and falls back on hysteresis, so the app never sheds load to save a few dollars. The **\$3,000 native-apps recipe** delivers the same everything-on + all-platforms + standing-auto- scaling + load-test posture for **~\$1,540-\$2,940** (Capacitor + Codemagic free-tier native builds + fixed-bid deploy); see the floor table below.

Older single figures in these docs predate the kit/automation + cost floor that trimmed the numbers; the banner above is current.

Posture: everything is ON, up, and running from the get-go. There is no build phase — the product ships feature-complete with every flag ON by default and pre-warms its own content, so launch is deploy + test + submit only. That is *why* these numbers are this low. See `EVERYTHING-ON-FROM-DAY-ONE.md`.

Prepared 2026-07-30 • Supersedes the 2026-07-28 estimate • Figures are planning targets, not quotes

This revision folds in the work completed tonight and re-states the posture you asked for: **everything running and on from the get-go**, with the developer cost driven to its honest floor. Tonight's additions are the auto-qualify → one-tap Premium path and the Services marketplace section brought to full parity with the other sections (serverless-GPU category tiles + subsections + search). Both shipped **prebuilt and switched ON by default**, so they add capability without adding launch cost.

Does tonight's work actually run?

Straight answer: it is **verified-correct, not yet runtime-proven** — and that distinction matters, so here is exactly what was and wasn't checked.

What passed, tonight, on every file: the structural audit (657 backend functions, 0 guardrail warnings), a TypeScript syntax sweep, brace/paren/bracket balance, JSON validity for the manifest and scheduler, every relative import resolving to a real file on disk, all four new backend functions registered in the manifest, and the frontend wiring (the Premium banner mounted in the dashboard, the Services page calling the GPU-category API, the page routed in `pages.config.js`). Each new function is also a near-line-for-line mirror of an existing, already-working one — `premiumAcceptOffer` mirrors `loyaltyEnroll`, `serviceStoreCategories` mirrors `appStoreCategories`, `aiServiceCategoryImages` mirrors `aiAppCategoryImages` — so it inherits proven patterns rather than inventing new ones.

What could **not** be checked here: a live boot. This build environment has no Deno runtime and no live Postgres, so nothing was executed against a real database. That first true runtime happens the moment you push to GitHub and Railway rebuilds. The one-command `e2e-smoke` in the kit is what proves the earn→unlock→purchase loop on the live instance, and the Premium banner only appears for a user who has actually hit the survey-day milestone, so plan to confirm both against a seeded test account after the first deploy. Nothing tonight touched money math or the payout rails, which keeps the runtime risk low.

Everything on from the get-go — the flag posture

The platform ships with essentially every capability **live at launch**, not dark. That is a code fact, not a promise: the built-in flag defaults turn the features on. The only things held off are held off for a specific external prerequisite — a processor account, an opt-in record, or a licensing question — not to save money, and each is a one-line flip once its prerequisite exists.

On by default (runs day one)	Off by default (and the specific reason)
Premium PPC · Loyalty/Rewards · auto-qualify Premium	Card charging — needs a live Stripe/PayPal processor + legal sign-off
Marketplace: Physical · Digital · Services (GPU tiles)	Affirm BNPL — needs Affirm merchant keys (real shippable goods only)
AI catalog, assistant, optimizer, self-learning, live experiments	SMS marketing — needs verifiable TCPA opt-in
AI concierge funnel + auto-launch on business pages (<code>ai_funnel</code>)	Teen accounts — needs verifiable parental consent + counsel
Results/benchmark auto-compilers + per-product stats	Store-credit purchase & P2P transfers — money-transmission risk
Tier 2 “Scale” pay-as-you-go ladder + deliverables (not credit)	Earnings projections — FTC earnings-claims review
Founding rollover/upgrade discount + \$2,000 premium gift boost	flexpay (installment credit) — licensed provider or attorney-confirmed self-financed
KYC survey · welcome rewards · Daily/Points Boost · layaway · purchase-payback	tier1_financed (recourse credit) — licensed creditor + counsel

Referrals · jackpots · social posting · consent-gated email re-engagement · telemetry · UX heatmap	goods_advance (advance credit) — licensed provider + counsel
Partner cash-out (closed-loop for regular users, ON)	

So “everything on” is already the shipped state — including every feature added this session. The right-hand column is unchanged in spirit: guardrails held for a specific external prerequisite (a processor, an opt-in, or a licensing/counsel question), each a one-line flag flip once that prerequisite lands. The three **credit** products added this session sit there too — off by design, **\$0 legal cost until you choose to unlock one**. The results-claims story runs compliant with no counsel: a hypothetical “how it works” shows until real data passes the sample threshold, then the substantiated figure auto-publishes with its basis.

One-time development — driven to the floor

Billed at \$75/hour. Because tonight’s features (and the AI, catalog, and marketplace layers before them) are prebuilt, the developer spends **zero** hours building product — only deploying, testing, and submitting. That is why the number keeps falling: there is less and less hand-work left.

Work item	Hours	Note
Accounts & API keys	0 billable	Owner-completed from the keys worksheet
Deploy backend + Postgres + frontend	5-7	Auto-migrating, same-origin frontend
Stand up the scheduler service (daily jobs)	1	Deploy <code>backend/scheduler/main.ts</code> as an always-on process (Handwork §E)
Pre-deploy validation	1	Scripted (<code>validate.sh</code>)
Configure AI / catalog / flip the cost floor on	1-2	Environment + flags only — no code, all prebuilt
Survey / earn-loop live test	3-5	<code>seed-demo</code> + <code>e2e-smoke</code>
QA pass	5-7	Against the QA test plan
Android submission (signed .aab)	8-11	Android submission kit + fastlane
Subtotal — PWA + Android	~24-34	≈ \$1,800-\$2,550
iOS submission (cloud CI, no Mac)	+4-8	iOS-without-a-Mac kit
Total — PWA + Android + iOS	~28-42	≈ \$2,050-\$3,150

Card payments stay off at launch, so the 0-8 “payments go-live” hours from prior estimates are out of the launch number entirely — that work moves to whenever you turn card charging on.

Ongoing runtime (monthly) — floor posture

Item	Cost
Hosting — backend + Postgres + the always-on scheduler service	~\$10-35/mo
Catalog + category images, one-time all-in	~\$11-19 once
LLM text — catalog, assistant, concierge funnel,	Now ~ \$0-20/mo with <code>AI_FORCE_CHEAP_TIER</code> (all

compilers	calls on Llama-8B), under the daily cap
AI / speech-to-text / TTS / email	\$0 (free tiers, or your self-hosted endpoints)
Legal & compliance	\$0 at launch — credit features off; optional one-time attorney read only if you unlock <code>self_financed</code> <code>flex</code> on day one
Geo-IP + exchange rates	Free tier
Realistic month one	~\$15-75 , capped

The one genuinely new recurring line is the **scheduler** — a second always-on process that runs the daily jobs (catalog seed, optimizer, self-learning, funnel benchmark compile, re-engagement sweep, product-stats compile, loyalty/points reconcile). It's a small service (a few \$/mo, or free if co-located), and without it the automation jobs don't fire — so it's now an explicit deploy step (Handwork \$E). Everything else this session added **zero** per-user cost: the concierge funnel, stats, and Tier 2 ladder are all reads over data you already store.

The floor — and why it isn't zero

Path	Year-one all-in	What it takes
Absolute shoestring (web PWA + Android, iOS as fast-follow)	~\$2,000-\$2,800	Dev \$1,725-2,475 · Play \$25 · domain ~\$15 · hosting ~\$120-180 · LLM capped ~\$60-180 · images ~\$0-19
Full three platforms (adds native iOS)	~\$2,900-\$4,000	+ iOS dev \$300-600 · Apple \$99/yr
Full three + AWS auto-scaling	~\$3,100-\$4,800	+ the auto-scaling infra floor (App Runner + single-AZ RDS trims it) — this was the old <i>uncapped</i> top end, before the codeable budget cap below
Full three + everything on + auto-scaling + load test (the complete launch, budget-capped)	~\$2,400-\$3,000	The full-scale ceiling brought down to \$3,000 by a codeable hard cap. The old \$4,800 top end assumed <i>uncapped</i> auto-scaling infra; a new monthly-budget lever (<code>INFRA_SCALE_MONTHLY_BUDGET_USD</code> , default \$48/mo) now derives the governor's normal max replicas from a dollar budget — <code>soft_max = min(max_instances, floor(budget / cost_per_instance))</code> — so steady-state scaling infra is ≤ 4 replicas $\approx$ \$48/mo $\approx$ \$576/yr , provably. “Everything on” still adds $\approx$ \$0 (AI/LLM on free Llama-8B, closed-loop revenue features \$0 at the margin and <i>generating</i> money); the load test is $\sim$ \$10-50 owner-run; legal stays \$0 while credit/money products stay counsel-gated. A real surge isn't shut out: <code>INFRA_SCALE_BURST_ABOVE_BUDGET (ON)</code> lets the governor auto-burst past the budget up to an 8-replica emergency ceiling to stay up, then fall back on hysteresis — so burst spend applies only for the minutes a spike lasts, and the annual number holds at $\sim$ \$3,000.
★ \$3,000 native-apps recipe (native iOS +	~\$1,540-\$2,940	The way to hit the complete posture and still stay under

Android, everything on, standing auto-scaling ON, load test)	\$3,000. Same feature set and the same standing auto-scaling as the row above — the savings come from <i>how you build and host</i>, not from turning anything off: (1) native iOS + Android via the Capacitor wrappers + Codemagic free-tier CI already in the repo (both store binaries built in the cloud, no Mac, no native rebuild) trims dev to ~\$1,300–\$2,100; (2) a fixed-bid deploy (owner does the store submissions from the kit) caps labor instead of billing hourly; (3) standing auto-scaling is ON from day one (governor <code>AUTO_SCALE_ENABLED=1</code>, <code>INFRA_SCALE_PROVIDER=railway</code>, min 2 instances) but costs only ~\$240–\$840/yr because Railway replicas are cheap and the governor scales <i>down</i> on hysteresis when load subsides; (4) load test stays ~\$10–50. Legal still \$0 while credit/money products stay counsel-gated. Net: ~\$1,540–\$2,940 year-one all-in — the everything-on + all-platforms + auto-scaling + load-test posture, delivered for ≤ \$3,000.
--	---

The hard external floor you cannot code away is about **\$40** on the shoestring path — Google Play \$25 plus a ~\$15 domain — because the PWA installs on iOS Safari, so iPhone users are covered without the \$99 Apple fee until you choose to ship the native iOS app. Add native iOS and the fixed floor becomes ~\$139 (Apple \$99 + Play \$25 + domain \$15). Below that sits only the human/external time that no automation removes: you creating the provider accounts, the first-run dashboard setup, and Google/Apple review. That residue — not code — is the real floor.

Bottom line

With this session’s additions — the AI concierge funnel, the results/benchmark/product-stats compilers, the Tier 2 “Scale” ladder and its deliverables, and the founding rollover/discount — all prebuilt and **on by default**, the launch cost did **not** rise. More of the platform is build-complete, so the developer still only deploys, tests, submits, and now also stands up the scheduler. A web + Android launch with **everything running from day one** lands at roughly **\$2,000–\$2,800 all-in for year one**; adding native iOS puts it near **\$2,900–\$4,000**. The irreducible fixed floor is **~\$40** shoestring (Play + domain), **~\$139** with native iOS. Recurring cost fell at the margins: the stronger cost floor (`AI_FORCE_CHEAP_TIER`) trims the LLM line to ~\$0–20/mo, and legal cost is **\$0** at launch because every credit product ships off behind its counsel gate. The only real risk to the number remains an iOS App Store rejection round — mitigated by the demo login and merit-not-gambling framing already in the kit.