

PROVISIONAL PATENT APPLICATION (DRAFT FOR COUNSEL)

Title of Invention: Integrated Closed-Loop Earn-to-Shop Marketplace with Graduated-Autonomy AI Automation and an Operator-Owned, Accuracy-Gated Custom AI Model Across Web and Mobile Surfaces

⚠ **DRAFT — NOT LEGAL ADVICE. NOT YET FILED.** This is an inventor's draft of a U.S. provisional patent application specification, assembled from the project's technical disclosures for a registered patent attorney to review, tighten, and file. It has **not** been reviewed for patentability (35 U.S.C. §§101/102/103), enablement, or written-description sufficiency. Do not rely on it as filed protection. **Section 13 ("What counsel must complete") lists the remaining work only a patent attorney should do.** Nothing herein should be publicly disclosed, demonstrated, sold, or offered for sale before counsel advises on the U.S. one-year statutory bar and on foreign absolute-novelty requirements. A provisional application is **not examined and is never "granted"**; it secures a priority date for 12 months, within which one or more non-provisional applications claiming priority to it must be filed.

0. Applicant, inventorship, and cross-reference (counsel to finalize)

- **Named inventor (sole, to date):** Ben Vick — benjaminjohnvick@gmail.com.
 - **Applicant / owner:** [Counsel to confirm — individual vs. an assignee entity (e.g., an LLC); record any assignment.]
 - **Priority claim:** This is an original provisional; no earlier application is claimed. [If any earlier provisional exists, counsel to insert the cross-reference under 35 U.S.C. §119(e).]
 - **Evidence of conception / reduction to practice:** dated Git commit history of repository [benjaminjohnvick-cmyk/playearningnexus](#); the full specification set in the Lawyer Packet; running code.
 - **Related disclosures incorporated by reference:** [PATENT-GROUNDWORK-AND-INVENTION-DISCLOSURE.md](#), [PATENT-FEATURE-AND-FUNCTION-DISCLOSURE-BRIEF.md](#), [CUSTOM-AI-MODEL-DESIGN-AND-ROADMAP.md](#), [WEBSITE-AND-APP-LOAD-SPEED.md](#), [EARN-WHILE-LOADING-SURVEY.md](#).
-

1. Technical field

Computer-implemented systems and methods for an integrated, closed-loop "earn-to-shop" digital marketplace operating across a responsive website and companion mobile applications, comprising: a consumer rewards economy funded by advertisers rather than by consumer markup; an artificial-intelligence advertising engine that generates, tests, and optimizes creative across owned and social surfaces; a platform-wide graduated-autonomy automation framework under which AI progressively assumes operational decisions subject to measured trust and permanent human/compliance gates on money, identity, and legal actions; and an operator-owned custom AI model, trained on first-party operational data, that is autonomously promoted to serve platform decisions only when its output accuracy exceeds an incumbent model on a per-function and aggregate basis.

2. Background

Existing get-paid-to, survey, loyalty, affiliate, and e-commerce platforms exhibit recurring technical and operational problems: rewards funded by consumer markups or opaque data resale; advertising creative produced and optimized manually and slowly; automation that is all-or-nothing (humans in every loop, or unguarded AI control); earning/referral/endorsement incentives that create disclosure, pyramid-scheme, money-transmission, and tax-reporting exposure bolted on after the fact rather than enforced by software; operating costs (AI inference, media rendering, storage) that scale unfavorably; and model-deployment tooling that selects or swaps machine-learning models on a single global metric or traffic split, without per-function safety gating or binding the swap to compliance controls. The present invention addresses each with specific technical mechanisms, several believed novel individually and non-obvious in combination.

3. Summary of the invention

The invention is a single integrated system of cooperating subsystems, any of which may be claimed independently and which are also claimed in combination:

1. A **closed-loop rewards economy** using non-cashable closed-loop store credit and an **earn-to-unlock** progression with an **earn-parity** mechanism keeping earned and paid paths equivalent, plus automated solvency reservation.
2. An **advertiser-funded model with a delivered-advertising-value guarantee** that measures and substantiates

delivered value while never promising revenue/ROI, enforced by an automated claim-compliance screen.

3. An **AI generative-creative engine** (including social video) that samples and predictively pre-scores candidates over a combinatorial concept space, renders only survivors, tests them on platform-owned surfaces, and tailors output toward measured best-responders, grounded in live trend/current-event data.
4. A **concept-polling loop** (head-to-head / MaxDiff) feeding the creative engine's playbook.
5. A **self-graduating video/creative autopilot** running the loop end-to-end behind a human approval gate that graduates to autonomy as trust is earned.
6. A generalized **Autonomy Kernel**: per-domain trust graduation (manual → earned → full) from measured human-agreement signals, with permanent money/identity/legal gates, a global kill switch, and budget caps.
7. **Automatic multi-surface feedback auto-collection** feeding the learning loops.
8. **Unbiased ("Fair Choice") choice presentation** yielding bias-controlled preference data.
9. A **self-learning creative playbook** with per-attribute, sample-smoothed scoring conditioning generation.
10. A **paid-endorser social program** with AI personalization, enforced disclosure, autonomy-gated auto-posting, and conversion-based rewards.
11. A **clawback-gated two-tier referral settlement** funded only from cleared, retained revenue.
12. A **whole-stack one-click cost-floor optimizer** with feature-preserving provider offload and automatic fallback.
13. A **compliance-as-code enforcement layer** binding jurisdiction/age/disclosure/money/consent/tax controls to the same switches that govern autonomy.
14. An **operator-owned custom AI model for e-commerce** trained from first-party operational data (incumbent labels + human approve/reject + optimizer outcomes), shadow-evaluated against an incumbent model's output accuracy **per function and in aggregate**, and autonomously promoted to serve **only when it exceeds the incumbent on every function individually and overall** (sustained), with automatic regression rollback, all behind permanent money/identity/legal human gates.
15. A **single operator-owned model operating an entire consumer website AND its companion mobile apps** from one unified cross-surface first-party corpus, with per-surface-and-function accuracy-gated switchover.
16. A **predictive "earn-while-it-loads" mechanism** that, upon predicting a load will exceed a human-perception latency budget, presents profiling questions and awards capped closed-loop credit until the load completes, while concurrently and budget-awarely preloading the application.

4. System architecture (implementation)

- **Front end:** responsive single-page web application (React/Vite) with a configuration-driven page router (~278 route-level pages), packaged as native/wrapped mobile applications (iOS/Android) sharing one code base and one service-worker-cached application shell.
- **Back end:** ~1,028 serverless functions (Deno runtime) dispatched via a manifest, plus ~225 typed subsystem modules; an in-process invocation bus composes functions without HTTP.
- **Data:** a document-style store (JSONB rows) with generated relational schema and expression indexes; an append-only money ledger and an append-only consent ledger; optimization-signal and learning-memory stores; ~369 persisted entity types.
- **Scheduling:** a cron-style scheduler drives recurring autonomous jobs.
- **AI providers:** provider-abstracted LLM/image/TTS/STT inference with runtime selection and automatic fallback; a single pluggable prediction interface (`serve(domain, context)`) with a backend switch selecting which model answers.
- **Money rails:** closed-loop store credit/points as the default settlement unit; external payout rails behind permanent gates, KYC, and a 1099 tax-export pipeline.

5. Brief description of the drawings (figures for counsel's draftsman to render)

Provisional applications may include informal drawings; the following figures are described so a patent draftsman can formalize them. Each reference numeral is a placeholder for counsel to normalize.

- **FIG. 1** — System overview: web + mobile clients, serverless function bus, data stores (money ledger, consent ledger, optimization signals), AI provider abstraction, and the model backend switch.
- **FIG. 2** — Autonomy Kernel state machine: per-domain graduation manual → earned → full; permanent-gate domains; kill switch; budget caps; the trust-signal inputs (approved runs, agreement rate, data depth).
- **FIG. 3** — Closed-loop economy: earn and paid paths, earn-parity equivalence, solvency reservation, ledger.
- **FIG. 4** — Generative-creative pipeline: concept-space sampling → predictive pre-scoring → render survivors → owned-surface testing → measurement → playbook update; concept-polling feedback.
- **FIG. 5** — Paid-endorser flow: personalization → enforced disclosure pinning → triple-gated auto-post → conversion measurement → reward settlement.

- **FIG. 6** — Custom-model lifecycle: first-party example assembly → shadow prediction → per-function + aggregate accuracy comparison vs. incumbent → sustained-streak promotion → serving → automatic regression rollback, with permanent gates intercepting money/identity/legal decisions regardless of backend.
- **FIG. 7** — Cross-surface single-model deployment across web and mobile apps from one unified corpus.
- **FIG. 8** — Earn-while-loading: load-latency prediction → question presentation + capped reward → concurrent budget-aware full-app preload → dismissal on load completion.
- **FIG. 9** — Compliance-as-code layer binding jurisdiction/age/disclosure/money/consent/tax controls to the autonomy switches.
- **FIG. 10** — Whole-stack cost-floor optimizer: lever set → provider offload → automatic fallback, features preserved.

6. Detailed description

The following incorporates, and is consistent with, the subsystem descriptions in [PATENT-GROUNDWORK-AND-INVENTION-DISCLOSURE.md](#) §5.1–5.16. Each subsystem is described functionally (for breadth) and with a representative implementation (for enablement and §101 specificity). Module names identify the technical reviewer's reference implementation; claims are not limited to the named modules.

6.1 Closed-loop earn-to-unlock economy with earn parity

A settlement unit of non-cashable, closed-loop store credit is the default medium of exchange. Users unlock premium capabilities either by paying or by platform activity, with an **earn-parity** function maintaining equivalence between the earned and paid paths so neither is advantaged; an automated **solvency reservation** reserves backing for outstanding credit before it is spendable. Append-only ledger entries make balances auditable. *Mechanism*: money primitives ([balance.ts](#), append-only ledger), parity computation, reservation on issuance. *Breadth*: claim the parity-maintained dual-path unlock + solvency reservation independent of the specific credit unit.

6.2 Advertiser-funded model with delivered-value guarantee and automated claim screening

Revenue derives from advertisers, not consumer markup. A **delivered-advertising-value** computation substantiates value delivered (e.g., impressions × conventional rate, measured engagement) while an automated screen **blocks any claim that promises revenue/ROI**, enforcing the guarantee's "value delivered, not results promised" posture in software. *Breadth*: claim the software-enforced separation of substantiated delivered-value from prohibited results-promises.

6.3 Generative-creative optimization over a combinatorial space with owned-surface testing

A very large combinatorial creative space is **sampled and predictively pre-scored**; only predicted winners are **rendered** (saving inference/render cost); survivors are tested on **platform-owned surfaces**; quantifiable response is measured; generation is tailored toward best responders; and generation is **grounded in live trend/current-event data**. *Breadth*: claim the sample → predict → render-only-winners → owned-surface-measure → self-tailor loop, independent of medium (video/image/copy).

6.4 Concept-polling loop

AI-generated concepts are turned into head-to-head / MaxDiff user polls whose results condition subsequent generation. *Breadth*: claim preference-poll-conditioned generative advertising.

6.5 Self-graduating creative/video autopilot

The end-to-end loop runs behind a human approval gate that **self-graduates** to autonomy on measured trust. (Implements §6.6 for the creative domain.)

6.6 The Autonomy Kernel — graduated-trust decision automation with permanent compliance gates

Every automatable decision is classified into a **domain**. Each domain graduates **manual → earned → full** based on measured trust signals (count of approved runs, human-agreement rate, data depth) against thresholds; reaching "earned" permits auto-approval within caps. **Permanent-gate** domains (money, identity, legal, high-risk) can **never** auto-approve regardless of trust. A **global kill switch** forces all domains to manual; **budget caps** bound autonomous spend. *Mechanism*: [autonomy-kernel.ts](#) (policy math), [autonomyDecide](#) (records an [AutonomyDecision](#), computes trust, returns auto-approve vs. await-human). *Breadth* — *lead claim*: a reusable controller that graduates per-domain autonomy from measured human-agreement while holding a fixed set of domains permanently gated, with kill switch and caps.

6.7 Automatic multi-surface feedback auto-collection

Implicit signals are harvested from customer-facing surfaces (navigation, dwell, interaction, outcome) and routed to the

learning loops without requiring explicit user input. *Breadth*: claim automatic cross-surface implicit- signal harvesting driving autonomy/creative learning.

6.8 Unbiased "Fair Choice" presentation

Choices (topics/ads) are presented **without favoring any option** (order/prominence controlled) so that selection data is bias-controlled. *Breadth*: claim bias-controlled choice elicitation for passive preference data.

6.9 Self-learning creative playbook

Each creative is tagged across attribute axes; per-attribute win rates are learned from measured outcomes with **sample smoothing** (shrinkage toward a prior for low-sample attributes) and condition future generation across owned and endorser channels. *Breadth*: claim the shared, sample-smoothed, cross-channel attribute playbook.

6.10 Paid-endorser social program

Opted-in members' own accounts post AI-**personalized, disclosure-enforced** advertisements; disclosure (e.g.,

ad) is pinned at generation and unremovable; posting is gated by the Autonomy Kernel; rewards are a share

of **measured conversion value**; copy self-improves on conversion data. *Breadth*: claim enforced-disclosure, approved-claim-pinned endorsement distribution with autonomy-graduated auto-posting and measured-conversion rewards.

6.11 Clawback-gated two-tier referral settlement

Referral incentives pay closed-loop credit; an advertiser bonus is released **only after the referred advertiser's payment clears and a hold window elapses** (clawback-gated), funded from cleared, retained revenue. *Breadth*: claim clawback-gated settlement from cleared revenue (single-tier; explicitly not a downline).

6.12 Whole-stack one-click cost-floor optimizer

A single control pulls every configuration lever to move inference/media/voice/storage onto free or cheaper providers, **preserving every feature**, with **automatic fallback** on provider failure. *Breadth*: claim the feature-preserving, fallback-protected whole-stack cost-floor controller.

6.13 Compliance-as-code enforcement layer

Jurisdiction and age gating at the point of action; disclosure enforcement at content generation; closed-loop money primitives with an append-only auditable ledger; a consent ledger; KYC and a filing-ready 1099 export; reversible feature flags that ship sensitive features **off** behind named switches. *Breadth*: claim the integrated guardrails-in-code layer binding money/identity/disclosure/jurisdiction/consent/tax controls to the same switches that govern autonomy (\$6.6).

6.14 Operator-owned custom AI model for e-commerce with function-by-function, accuracy-gated autonomous switchover

An incumbent general model initially operates the platform and serves two roles: it **labels** the operator's first-party data as it operates (its proposed action per e-commerce function plus the human approve/reject on it) and it is the **accuracy benchmark**. A pipeline assembles those labels, optimizer win/loss outcomes, and per-feature quality into **first-party, PII-minimized** (*context* → *label/reward*) training examples. The operator's model runs **in shadow from day one** behind a single pluggable interface, producing a candidate for every decision. An evaluator scores the candidate's **output accuracy against the incumbent's**, on the same human ground truth, **both per e-commerce function (domain) AND in aggregate**. The serving backend is **automatically promoted** to the operator's model **only when it exceeds the incumbent on EVERY function individually by a margin AND in aggregate**, each over a minimum sample count and **sustained across consecutive evaluations**, and is **automatically rolled back** on regression. Money/identity/legal decisions remain **permanently human-gated** regardless of which model serves; the training-data export is gated behind a counsel-controlled switch. *Mechanism*: [custom-model.ts](#) (interface + backend switch + built-in data-learner), [model-eval.ts](#) (per-function + aggregate gate, auto-promote/rollback), [model-training.ts](#) (first-party example assembly). *Breadth* — *key claim*: promotion gated on per-function **and** aggregate out-performance of an incumbent, with the incumbent as simultaneous labeler and live benchmark while the challenger shadows, sustained-streak promotion, automatic regression rollback, and binding to permanent compliance gates.

6.15 Single operator-owned model operating an entire website AND its mobile apps from one cross-surface corpus

One operator-owned model, trained on a **unified first-party corpus spanning both the website and the native mobile apps** (same account, same closed-loop economy, same decision domains), operates the **whole end-to-end consumer experience** across surfaces, making **consistent decisions** rather than a separate model/ruleset per platform, under the same accuracy-gated switchover applied **per-surface-and-function**, with permanent gates on every surface. *Breadth*: claim the single cross-surface model operating an entire consumer property end-to-end from one unified corpus, with cross-surface consistency as an objective.

6.16 Predictive earn-while-it-loads with concurrent budget-aware preload

Upon **predicting in advance** (from measured real-user latency percentiles, connection quality, and degraded- state signals) that a load will exceed a **human-perception latency budget** (e.g., below the ~100 ms threshold at which a response is perceived as instantaneous), the system presents generic, non-sensitive profiling questions and awards **capped, closed-loop store credit** per answer until the load completes; **concurrently**, it **preloads the remaining application** in a **budget-aware** manner (idle-scheduled, one unit at a time, self-throttling on measured congestion) so subsequent navigation is instantaneous, without exceeding the foreground latency budget. *Breadth*: claim the predict-slow-load → reward-bearing interstitial + concurrent budget-aware preload combination.

6.17 Supporting subsystems (breadth of the system)

Surveys/evidence; gamification; marketplace/catalog and one-click member storefronts; advertiser tiers and financing; payouts and fraud scoring; household/teen accounts; AI support/dispute automation; localization; white-label/tenancy. These establish the full system scope for the combination ("system") claims.

7. Claims (claim seeds — broad independent claims with dependent ladders; counsel to finalize)

Provisional applications need not contain formal claims, but these seeds document the intended scope and support the priority date. They are drafted broad with narrower fallbacks. Counsel to re-draft for §101 eligibility, antecedent basis, single-sentence form, and to decide which become independent applications.

Independent Claim A (integrated system — broadest). A computer-implemented system comprising a web client and a mobile-application client sharing a common application shell; a serverless back end; a closed-loop settlement unit; and an autonomy controller that classifies operational decisions into domains and graduates each domain's automation from manual to autonomous based on measured human-agreement signals while holding a predefined set of money, identity, and legal domains permanently subject to human approval. - *Dep. A1*: wherein a global control forces all domains to manual (kill switch). - *Dep. A2*: wherein autonomous actions are bounded by configurable budget caps. - *Dep. A3*: further comprising an append-only money ledger and an append-only consent ledger. - *Dep. A4*: wherein rewards are funded by advertiser contributions rather than consumer markup. - *Dep. A5*: further comprising the custom-model subsystem of Claim E.

Independent Claim B (Autonomy Kernel — method). A method of automating operational decisions comprising: assigning each decision to a domain; computing, per domain, a trust measure from a count of human-approved actions, a human-agreement rate, and a data-depth measure; permitting autonomous execution of a domain's actions only when the trust measure satisfies domain thresholds; and unconditionally requiring human approval for a predefined set of money/identity/legal domains irrespective of the trust measure. - *Dep. B1-B4*: kill switch; budget caps; "clean-approval" (unmodified-approval) weighting of the agreement rate; per-domain threshold overrides.

Independent Claim C (generative-creative optimization). A method comprising defining a combinatorial creative space; predictively pre-scoring candidate creatives; rendering only candidates exceeding a score threshold; publishing rendered candidates to operator-owned surfaces; measuring response; and conditioning subsequent generation on measured best-responders; optionally grounded in live external trend data. - *Dep. C1-C4*: concept-polling conditioning; sample-smoothed per-attribute playbook; cross-channel playbook sharing; render-cost-avoidance limitation.

Independent Claim D (enforced-disclosure endorsement distribution). A method of distributing advertisements through opted-in members' own social accounts wherein a required disclosure is pinned at content generation and not removable, posting is gated by the autonomy controller of Claim B, and rewards are computed from measured conversion value.

Independent Claim E (operator-owned accuracy-gated custom model — key standalone). A computer-implemented method comprising: operating a platform with an incumbent model; recording, as training examples, the incumbent's per-function proposed decisions together with human approve/reject labels; training a challenger model on said first-party examples; executing the challenger in shadow to produce candidate decisions; measuring output accuracy of the challenger and of the incumbent against the human labels **both per function and in aggregate**; switching a serving backend to the challenger **only when the challenger's accuracy exceeds the incumbent's on every function individually by at least a margin and in aggregate, sustained across a plurality of successive**

evaluations; automatically reverting the serving backend upon a subsequent regression; and executing a predefined set of money/identity/legal decisions under mandatory human approval regardless of the serving backend. - *Dep. E1*: wherein the incumbent model both labels the training data and serves as the live accuracy benchmark. - *Dep. E2*: wherein a function is eligible to gate the switch only after accumulating a minimum number of evaluated decisions. - *Dep. E3*: wherein training examples are PII-minimized. - *Dep. E4*: wherein the challenger is served via a pluggable interface selectable by configuration without code change. - *Dep. E5 (cross-surface)*: wherein a single challenger model operates both a website and a mobile application from a unified cross-surface corpus, and the switching is performed per-surface-and-function. - *Dep. E6*: wherein the training-data export is disabled behind a controlled switch.

Independent Claim F (predictive earn-while-loading + preload). A method comprising: predicting, from measured latency percentiles and connection signals, that a load will exceed a perception-latency budget; responsively presenting profiling prompts and awarding capped closed-loop credit per response until the load completes; and concurrently preloading remaining application code in an idle-scheduled, self-throttling manner that keeps foreground interaction latency within the budget.

Independent Claim G (clawback-gated referral settlement). A method of settling referral incentives in closed-loop credit wherein a referred-advertiser bonus is released only after the advertiser's payment clears and a hold window elapses, funded from cleared retained revenue.

Independent Claim H (whole-stack cost-floor controller). A method wherein a single control reconfigures inference, media, voice, and storage providers to a lowest-cost configuration while preserving enabled features, with automatic fallback to an alternate provider on failure.

Independent Claim I (compliance-as-code layer). A system binding jurisdiction gating, age gating, generation- time disclosure enforcement, append-only money and consent ledgers, a KYC/1099 pipeline, and reversible default-off feature flags to the same controls that govern the autonomy controller of Claim B.

(Counsel: consider one or two broad "system" claims (A, incorporating E) plus method families B-I, and evaluate which subjects warrant separate non-provisional filings versus dependent claims.)

8. Abstract

A cross-surface (web and mobile) closed-loop "earn-to-shop" marketplace in which rewards are advertiser-funded, AI generates and optimizes advertising creative on owned surfaces, and a graduated-autonomy controller lets AI assume operational decisions per domain under measured trust while permanently gating money/identity/legal decisions to human approval. An operator-owned AI model, trained on first-party operational data labeled by an incumbent model, runs in shadow and is automatically promoted to serve — and automatically rolled back — based on exceeding the incumbent's output accuracy on a per-function and aggregate basis across web and mobile, with compliance controls enforced in code.

9. Enablement pointers (for written-description support)

The reference implementation comprises ~1,028 serverless functions, ~225 subsystem modules, ~369 persisted entity types, and ~278 client pages, enumerated in [PATENT-FEATURE-AND-FUNCTION-DISCLOSURE-BRIEF.md](#) (incorporated by reference), providing written-description and enablement support for the claimed functions. Representative modules are named per subsystem in Section 6.

10. Distinctions over likely prior art (to aid counsel's search)

- vs. get-paid-to/survey apps: closed-loop non-cashable settlement + earn parity + advertiser-pool solvency reservation, not cash-out from data resale.
 - vs. ad-optimization / A-B tools: pre-production preference polling + predictive pre-scoring + render-only-winners
 - trend grounding as one loop, plus a shared cross-channel attribute playbook.
 - vs. RPA / workflow automation: autonomy that graduates from measured per-domain human agreement with permanent compliance gates, not static rules or blanket AI control.
 - vs. AutoML / champion-challenger / A-B model deployment: promotion gated on beating the incumbent **per function AND in aggregate** (not one global metric or a traffic-split winner); the incumbent is simultaneously labeler and live benchmark; promotion requires a sustained multi-evaluation streak with automatic regression rollback; and the autonomous swap is bound to permanent money/identity/legal human gates. Training corpus is the operator's own first-party, PII-minimized operational data.
 - vs. influencer platforms: disclosure enforced at generation and unremovable; rewards tied to measured conversion; posting gated by earned autonomy.
 - vs. MLM/referral software: single-tier, clawback-gated settlement from cleared revenue; not a downline.
-

11. Matters expressly reserved / flagged (do not treat as decided)

- Whether using an AI provider's outputs to label data for training a separate model is permitted under that provider's terms — a terms-of-service + privacy question (see [MODEL-TRAINING-DATA-COUNSEL-NOTE.md](#)).
 - The games-to-retail repositioning (2026-09-09): counsel may prefer claim language framed around retail/rewards with games as one searchable category.
 - Which elements to protect as **trade secret** rather than patent (exact thresholds, margins, scoring formulas), since a published patent teaches them.
-

12. What a provisional gives you, concretely

A filed provisional establishes a **priority date** for everything disclosed herein, permits "patent pending," is not examined, and lasts **12 months**, within which one or more **non-provisional** applications claiming priority must be filed or the priority is lost. Broad disclosure now is advantageous because **new matter cannot be added** later; narrowing happens in the non-provisionals.

13. WHAT COUNSEL MUST COMPLETE — the remaining 10-20% (only a patent attorney should do this)

This draft intentionally stops short of tasks that require a registered practitioner's judgment. Counsel should:

1. **Patentability & §101 eligibility.** Assess each subject against *Alice/Mayo*; re-draft claims as specific technical improvements (not abstract business methods); decide which subjects are eligible at all.
 2. **Prior-art search (§§102/103)** on each independent claim (especially E — model-deployment art is crowded) and set claim scope accordingly.
 3. **Formal claim drafting.** Convert the Section 7 seeds into proper single-sentence claims with correct antecedent basis, one independent claim per invention family, and properly nested dependents; decide the broad-vs-narrow ladder depth.
 4. **Filing strategy / restriction.** Decide whether to file **one** non-provisional with multiple inventions (and accept likely restriction requirements → divisionals) or **several** targeted non-provisionals off this one provisional priority date; plan the 12-month budget.
 5. **Drawings.** Have a draftsman render FIGS. 1-10 to USPTO standards and normalize reference numerals.
 6. **Inventorship & ownership.** Confirm sole vs. joint inventorship; execute any assignment to an entity; verify no third-party/contractor IP or open-source licensing taints the claimed matter.
 7. **Statutory-bar & foreign-filing timing.** Confirm no public disclosure/sale/demo has started the U.S. one-year clock; advise on PCT / foreign absolute-novelty before any launch; calendar the 12-month non-provisional deadline and any Paris Convention dates.
 8. **Trade-secret carve-out.** Decide what to **remove** from the public specification (exact thresholds, margins, formulas) and keep as trade secret.
 9. **Third-party-AI and ToS review.** Resolve the Section 11 reserved matters before claims rely on incumbent-model labeling.
 10. **Enablement/written-description sufficiency.** Confirm the specification (plus incorporated disclosures) enables each claim, and add any implementation detail an examiner would require.
 11. **Entity status & fees.** Determine micro/small/large entity status; file the provisional cover sheet (SB/16), specification, any drawings, and fee via Patent Center.
 12. **Title & claim-set finalization.** Approve or revise the title and the final independent-claim selection.
-

Prepared 2026-09-13 as an inventor's draft for patent counsel. Not legal advice; not filed; not a substitute for a registered patent attorney's review and filing.